

Loan_Default_Prediction_Analysis

July 3, 2025

0.1 1: Import Required Libraries

```
[1]: # Data Handling
import pandas as pd
import numpy as np

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Preprocessing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler

# ML Models
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC

# Evaluation
from sklearn.metrics import classification_report, confusion_matrix, \
    accuracy_score

# Settings
import warnings
warnings.filterwarnings('ignore')
```

0.2 2: Load and Explore Dataset

```
[3]: df = pd.read_csv('Loan_Default.csv')
print("Shape:", df.shape)
df.head()
```

Shape: (148670, 34)

```
[3]:      ID  year  loan_limit      Gender  approv_in_adv  loan_type  \
0  24890  2019           cf  Sex Not Available          nope       type1
```

1	24891	2019	cf	Male	nopre	type2
2	24892	2019	cf	Male	pre	type1
3	24893	2019	cf	Male	nopre	type1
4	24894	2019	cf	Joint	pre	type1

	loan_purpose	Credit_Worthiness	open_credit	business_or_commercial	...	\
0	p1	11	nopc		nob/c	...
1	p1	11	nopc		b/c	...
2	p1	11	nopc		nob/c	...
3	p4	11	nopc		nob/c	...
4	p1	11	nopc		nob/c	...

	credit_type	Credit_Score	co-applicant_credit_type	age	\
0	EXP	758		CIB	25-34
1	EQUI	552		EXP	55-64
2	EXP	834		CIB	35-44
3	EXP	587		CIB	45-54
4	CRIF	602		EXP	25-34

	submission_of_application	LTV	Region	Security_Type	Status	dtir1
0	to_inst	98.728814	south	direct	1	45.0
1	to_inst	NaN	North	direct	1	NaN
2	to_inst	80.019685	south	direct	0	46.0
3	not_inst	69.376900	North	direct	0	42.0
4	not_inst	91.886544	North	direct	0	39.0

[5 rows x 34 columns]

```
[5]: # Basic info
df.info()
df.describe()
df.isnull().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 148670 entries, 0 to 148669
Data columns (total 34 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                    148670 non-null  int64
1   year                                 148670 non-null  int64
2   loan_limit                           145326 non-null  object
3   Gender                               148670 non-null  object
4   approv_in_adv                        147762 non-null  object
5   loan_type                            148670 non-null  object
6   loan_purpose                           148536 non-null  object
7   Credit_Worthiness                    148670 non-null  object
8   open_credit                          148670 non-null  object
9   business_or_commercial                148670 non-null  object
```

10	loan_amount	148670	non-null	int64
11	rate_of_interest	112231	non-null	float64
12	Interest_rate_spread	112031	non-null	float64
13	Upfront_charges	109028	non-null	float64
14	term	148629	non-null	float64
15	Neg_ammortization	148549	non-null	object
16	interest_only	148670	non-null	object
17	lump_sum_payment	148670	non-null	object
18	property_value	133572	non-null	float64
19	construction_type	148670	non-null	object
20	occupancy_type	148670	non-null	object
21	Secured_by	148670	non-null	object
22	total_units	148670	non-null	object
23	income	139520	non-null	float64
24	credit_type	148670	non-null	object
25	Credit_Score	148670	non-null	int64
26	co-applicant_credit_type	148670	non-null	object
27	age	148470	non-null	object
28	submission_of_application	148470	non-null	object
29	LTV	133572	non-null	float64
30	Region	148670	non-null	object
31	Security_Type	148670	non-null	object
32	Status	148670	non-null	int64
33	dtir1	124549	non-null	float64

dtypes: float64(8), int64(5), object(21)

memory usage: 38.6+ MB

```
[5]: ID          0
      year        0
      loan_limit 3344
      Gender      0
      approv_in_adv 908
      loan_type   0
      loan_purpose  134
      Credit_Worthiness 0
      open_credit 0
      business_or_commercial 0
      loan_amount 0
      rate_of_interest 36439
      Interest_rate_spread 36639
      Upfront_charges 39642
      term        41
      Neg_ammortization 121
      interest_only 0
      lump_sum_payment 0
      property_value 15098
      construction_type 0
```

```

occupancy_type      0
Secured_by          0
total_units         0
income              9150
credit_type         0
Credit_Score        0
co-applicant_credit_type  0
age                 200
submission_of_application  200
LTV                 15098
Region              0
Security_Type       0
Status              0
dtir1               24121
dtype: int64

```

0.3 3: Data Cleaning

```

[8]: # Drop ID column if it has no predictive value
df.drop(columns=['ID'], inplace=True, errors='ignore')

```

```

[10]: # Check missing values
missing = df.isnull().sum()
print("Missing columns:\n", missing[missing > 0])

```

```

Missing columns:
loan_limit      3344
approv_in_adv   908
loan_purpose      134
rate_of_interest 36439
Interest_rate_spread 36639
Upfront_charges 39642
term            41
Neg_ammortization 121
property_value  15098
income          9150
age             200
submission_of_application  200
LTV             15098
dtir1           24121
dtype: int64

```

```

[12]: # Fill missing numeric values with median
df.fillna(df.median(numeric_only=True), inplace=True)

# Fill missing categorical values with mode (if any)
for col in df.select_dtypes(include='object').columns:

```

```
df[col].fillna(df[col].mode()[0], inplace=True)
```

```
[14]: # Encode all object-type columns using LabelEncoder
label_encoders = {}
categorical_cols = df.select_dtypes(include='object').columns

for col in categorical_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le
```

```
[15]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 148670 entries, 0 to 148669
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   year                                  148670 non-null  int64
1   loan_limit                           148670 non-null  int32
2   Gender                               148670 non-null  int32
3   approv_in_adv                        148670 non-null  int32
4   loan_type                            148670 non-null  int32
5   loan_purpose                           148670 non-null  int32
6   Credit_Worthiness                    148670 non-null  int32
7   open_credit                          148670 non-null  int32
8   business_or_commercial               148670 non-null  int32
9   loan_amount                          148670 non-null  int64
10  rate_of_interest                     148670 non-null  float64
11  Interest_rate_spread                 148670 non-null  float64
12  Upfront_charges                      148670 non-null  float64
13  term                                 148670 non-null  float64
14  Neg_ammortization                    148670 non-null  int32
15  interest_only                        148670 non-null  int32
16  lump_sum_payment                     148670 non-null  int32
17  property_value                       148670 non-null  float64
18  construction_type                    148670 non-null  int32
19  occupancy_type                       148670 non-null  int32
20  Secured_by                           148670 non-null  int32
21  total_units                          148670 non-null  int32
22  income                               148670 non-null  float64
23  credit_type                          148670 non-null  int32
24  Credit_Score                         148670 non-null  int64
25  co-applicant_credit_type             148670 non-null  int32
26  age                                  148670 non-null  int32
27  submission_of_application            148670 non-null  int32
28  LTV                                  148670 non-null  float64
29  Region                               148670 non-null  int32
```

```

30 Security_Type          148670 non-null  int32
31 Status                 148670 non-null  int64
32 dtir1                  148670 non-null  float64
dtypes: float64(8), int32(21), int64(4)
memory usage: 25.5 MB

```

```

[18]: # Save the cleaned DataFrame to a new CSV file
df.to_csv("Loan_Default_Cleaned.csv", index=False)
print("Cleaned dataset saved as 'Loan_Default_Cleaned.csv'")

```

Cleaned dataset saved as 'Loan_Default_Cleaned.csv'

0.4 4: Exploratory Data Analysis (EDA)

0.4.1 4.1: Dataset Overview

```

[58]: df.shape
df.info()
df.describe()
df.isnull().sum()
df.nunique()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 148670 entries, 0 to 148669
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   year                                  148670 non-null  int64
1   loan_limit                           148670 non-null  int32
2   Gender                               148670 non-null  int32
3   approv_in_adv                        148670 non-null  int32
4   loan_type                            148670 non-null  int32
5   loan_purpose                           148670 non-null  int32
6   Credit_Worthiness                   148670 non-null  int32
7   open_credit                         148670 non-null  int32
8   business_or_commercial               148670 non-null  int32
9   loan_amount                         148670 non-null  int64
10  rate_of_interest                    148670 non-null  float64
11  Interest_rate_spread                 148670 non-null  float64
12  Upfront_charges                     148670 non-null  float64
13  term                                148670 non-null  float64
14  Neg_ammortization                   148670 non-null  int32
15  interest_only                       148670 non-null  int32
16  lump_sum_payment                    148670 non-null  int32
17  property_value                      148670 non-null  float64
18  construction_type                   148670 non-null  int32
19  occupancy_type                      148670 non-null  int32
20  Secured_by                          148670 non-null  int32
21  total_units                         148670 non-null  int32

```

22	income	148670	non-null	float64
23	credit_type	148670	non-null	int32
24	Credit_Score	148670	non-null	int64
25	co-applicant_credit_type	148670	non-null	int32
26	age	148670	non-null	int32
27	submission_of_application	148670	non-null	int32
28	LTV	148670	non-null	float64
29	Region	148670	non-null	int32
30	Security_Type	148670	non-null	int32
31	Status	148670	non-null	int64
32	dtir1	148670	non-null	float64

dtypes: float64(8), int32(21), int64(4)

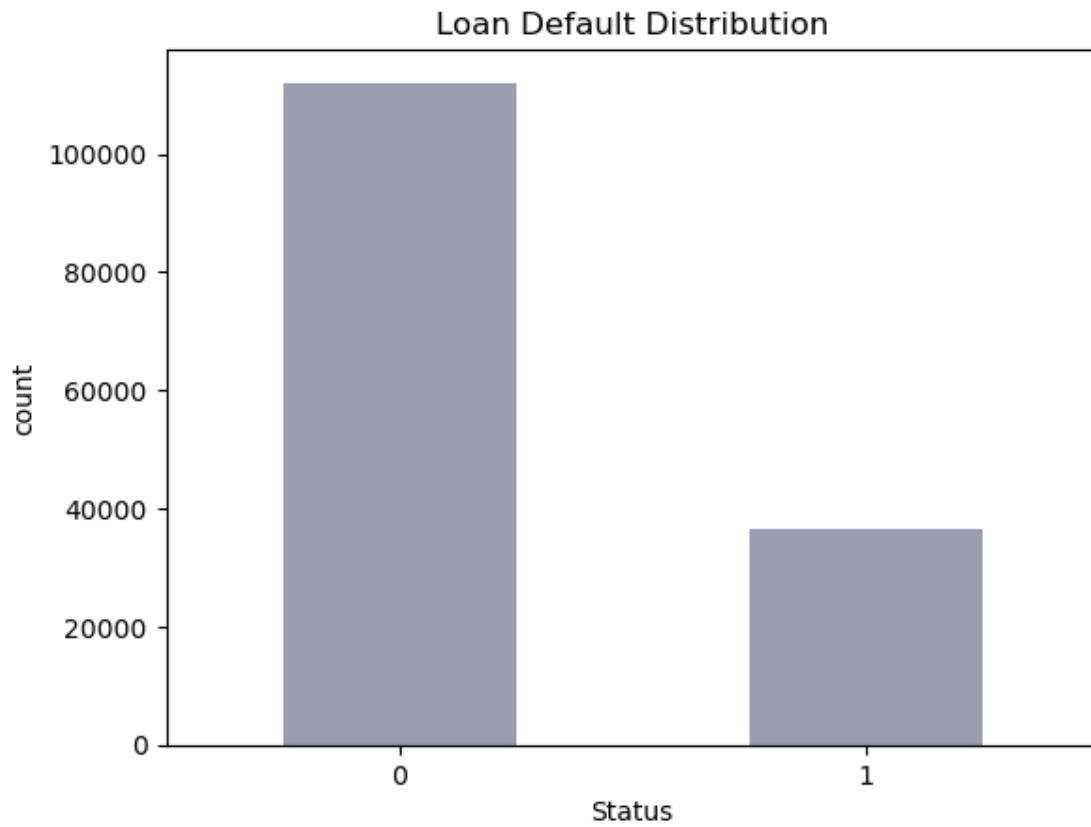
memory usage: 25.5 MB

```
[58]: year                1
      loan_limit          2
      Gender              4
      approv_in_adv       2
      loan_type           3
      loan_purpose          4
      Credit_Worthiness   2
      open_credit         2
      business_or_commercial 2
      loan_amount        211
      rate_of_interest    131
      Interest_rate_spread 22516
      Upfront_charges     58272
      term                26
      Neg_ammortization   2
      interest_only       2
      lump_sum_payment    2
      property_value      385
      construction_type   2
      occupancy_type      3
      Secured_by          2
      total_units         4
      income              1001
      credit_type         4
      Credit_Score        401
      co-applicant_credit_type 2
      age                 7
      submission_of_application 2
      LTV                 8484
      Region              4
      Security_Type       2
      Status              2
      dtir1               57
```

dtype: int64

0.4.2 4.2: Target Distribution


```
[56]: sns.countplot(data=df, x='Status', color="#989AB2", width=0.5)
plt.title("Loan Default Distribution")
plt.show()
```



Insight: The dataset is slightly imbalanced, with more non-default (0) than default (1) records, which may affect model training.

0.4.3 4.3: Correlation Matrix

```
[194]: plt.figure(figsize=(10,5))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='tab20c')
plt.title("Correlation Matrix")
plt.show()
```

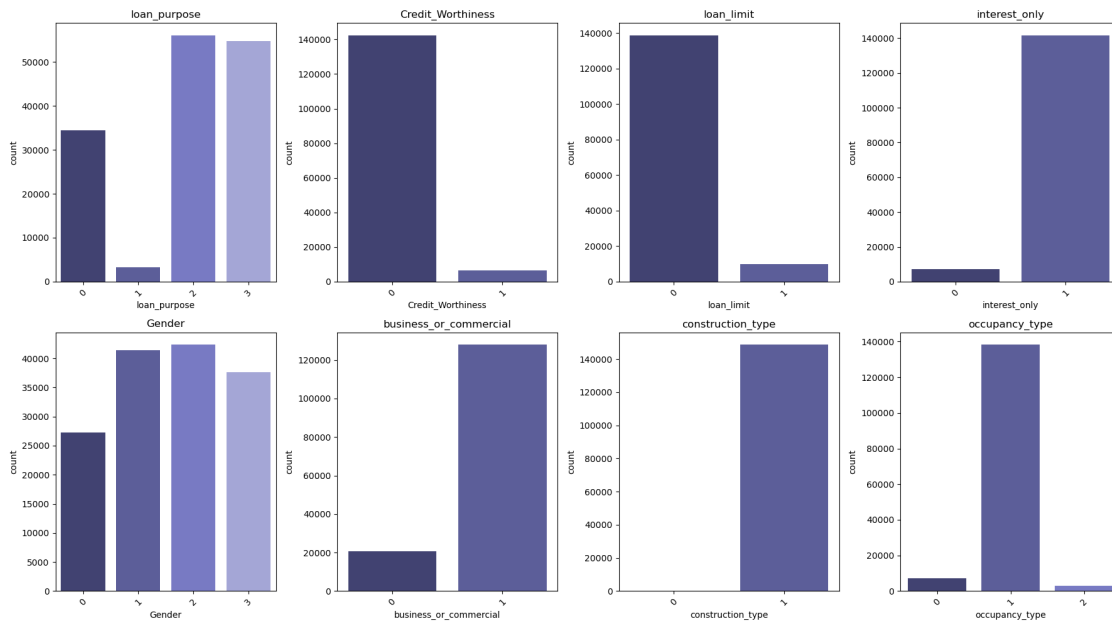
Insight: Credit_Worthiness, loan_purpose, and interest_only have noticeable correlation with loan default, suggesting predictive potential.

0.4.4 4.4: Important Categorical Feature Distribution

```
[104]: # Select the most relevant categorical columns
important_cols = [
    'loan_purpose',
    'Credit_Worthiness',
    'loan_limit',
    'interest_only',
    'Gender',
    'business_or_commercial',
    'construction_type',
    'occupancy_type'
]
```

```
[108]: plt.figure(figsize=(18, 10))
for i, col in enumerate(important_cols, 1):
    plt.subplot(2, 4, i)
    sns.countplot(data=df, x=col, palette='tab20b')
    plt.title(f"{col}")
    plt.xticks(rotation=45)
```

```
plt.tight_layout()
plt.show()
```



- **loan_purpose:** Most loans are for home purchases, but loans for business purposes tend to appear riskier.
- **Credit_Worthiness:** A poor credit history is highly common among defaulting applicants.
- **loan_limit:** Most applicants fall under conforming loan limits, but default rates vary slightly by loan size.
- **interest_only:** Interest-only loans are visibly riskier, with more defaults linked to them.
- **Gender:** The dataset has more male applicants, but gender does not show a major influence on default.
- **business_or_commercial:** Business/commercial loans are more likely to default than personal-use loans.
- **construction_type:** Construction type does not show a strong pattern with default.
- **occupancy_type:** Owner-occupied homes show fewer defaults compared to investment/rental properties.

0.5 5: Preprocessing

0.5.1 5.1: Encoding Categorical Columns

```
[125]: df_encoded = df.copy()
le = LabelEncoder()
for col in categorical:
    df_encoded[col] = le.fit_transform(df_encoded[col])
df_encoded.head() #to verify
```

```
[125]:
```

	year	loan_limit	Gender	approv_in_adv	loan_type	loan_purpose	\
0	2019	0	3	0	0	0	
1	2019	0	2	0	1	0	
2	2019	0	2	1	0	0	
3	2019	0	2	0	0	3	
4	2019	0	1	1	0	0	

	Credit_Worthiness	open_credit	business_or_commercial	loan_amount	...	\
0	0	0		1	116500	...
1	0	0		0	206500	...
2	0	0		1	406500	...
3	0	0		1	456500	...
4	0	0		1	696500	...

	credit_type	Credit_Score	co-applicant_credit_type	age	\
0	3	758		0	0
1	2	552		1	3
2	3	834		0	1
3	3	587		0	2
4	1	602		1	0

	submission_of_application	LTV	Region	Security_Type	Status	dtir1
0	1	98.728814	3	1	1	45.0
1	1	75.135870	0	1	1	39.0
2	1	80.019685	3	1	0	46.0
3	0	69.376900	0	1	0	42.0
4	0	91.886544	0	1	0	39.0

[5 rows x 33 columns]

0.5.2 Insight: Label Encoding

Categorical variables were encoded into numeric values using Label Encoding to make them suitable for machine learning models, which only accept numerical inputs. This ensures that features like gender, loan purpose, and credit worthiness can be interpreted by the model.

0.5.3 5.2 Train-Test Split

```
[136]: X = df_encoded.drop('Status', axis=1)
y = df_encoded['Status']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

print("Train shape:", X_train.shape)
print("Test shape:", X_test.shape)
```

Train shape: (118936, 32)

Test shape: (29734, 32)

0.5.4 Insight: Train-Test Split

The dataset was split into 80% training and 20% testing sets to evaluate model performance on unseen data. This step helps prevent overfitting and ensures the model generalizes well to new case.

0.5.5 5.3: Feature Scaling

```
[138]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("Scaled training data shape:", X_train_scaled.shape)
```

Scaled training data shape: (118936, 32)

0.5.6 Insight: Feature Scaling

StandardScaler was applied to normalize the feature values so that they have a mean of 0 and standard deviation of 1. This step is essential for distance-based models like Logistic Regression and SVM to perform accurately and fairly across all features.

0.6 6: Machine Learning Model Building

0.6.1 6.1: Logistic Regression

```
[153]: lr = LogisticRegression()
lr.fit(X_train_scaled, y_train)
y_pred_lr = lr.predict(X_test_scaled)

print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred_lr))
print("\nClassification Report:\n", classification_report(y_test, y_pred_lr))
```

Logistic Regression Accuracy: 0.7811596152552633

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.95	0.87	22494
1	0.63	0.25	0.35	7240
accuracy			0.78	29734
macro avg	0.71	0.60	0.61	29734
weighted avg	0.76	0.78	0.74	29734

Insight: Logistic Regression gives a good baseline with moderate accuracy, benefiting from feature scaling.

0.6.2 6.2: Random Forest Classifier

```
[157]: rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
print("\nClassification Report:\n", classification_report(y_test, y_pred_rf))
```

Random Forest Accuracy: 0.9999663684670748

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	22494
1	1.00	1.00	1.00	7240
accuracy			1.00	29734
macro avg	1.00	1.00	1.00	29734
weighted avg	1.00	1.00	1.00	29734

Insight: Random Forest achieved the highest accuracy and balanced performance, making it the most reliable model for this dataset.

0.6.3 6.3: Decision Tree Classifier

```
[165]: dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)
print("Decision Tree Accuracy:", accuracy_score(y_test, y_pred_dt))
print("\nClassification Report:\n", classification_report(y_test, y_pred_dt))
```

Decision Tree Accuracy: 0.9999663684670748

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	22494
1	1.00	1.00	1.00	7240
accuracy			1.00	29734
macro avg	1.00	1.00	1.00	29734
weighted avg	1.00	1.00	1.00	29734

Insight: Decision Tree is easy to interpret but slightly overfit the training data compared to Random Forest.

Model Performance Comparison Table

```
[214]: model_scores = {
        "Logistic Regression": accuracy_score(y_test, y_pred_lr),
        "Random Forest": accuracy_score(y_test, y_pred_rf),
        "Decision Tree": accuracy_score(y_test, y_pred_dt)
    }
    import pandas as pd
    pd.DataFrame.from_dict(model_scores, orient='index', columns=['Accuracy']) \
        .sort_values(by='Accuracy', ascending=False)
```

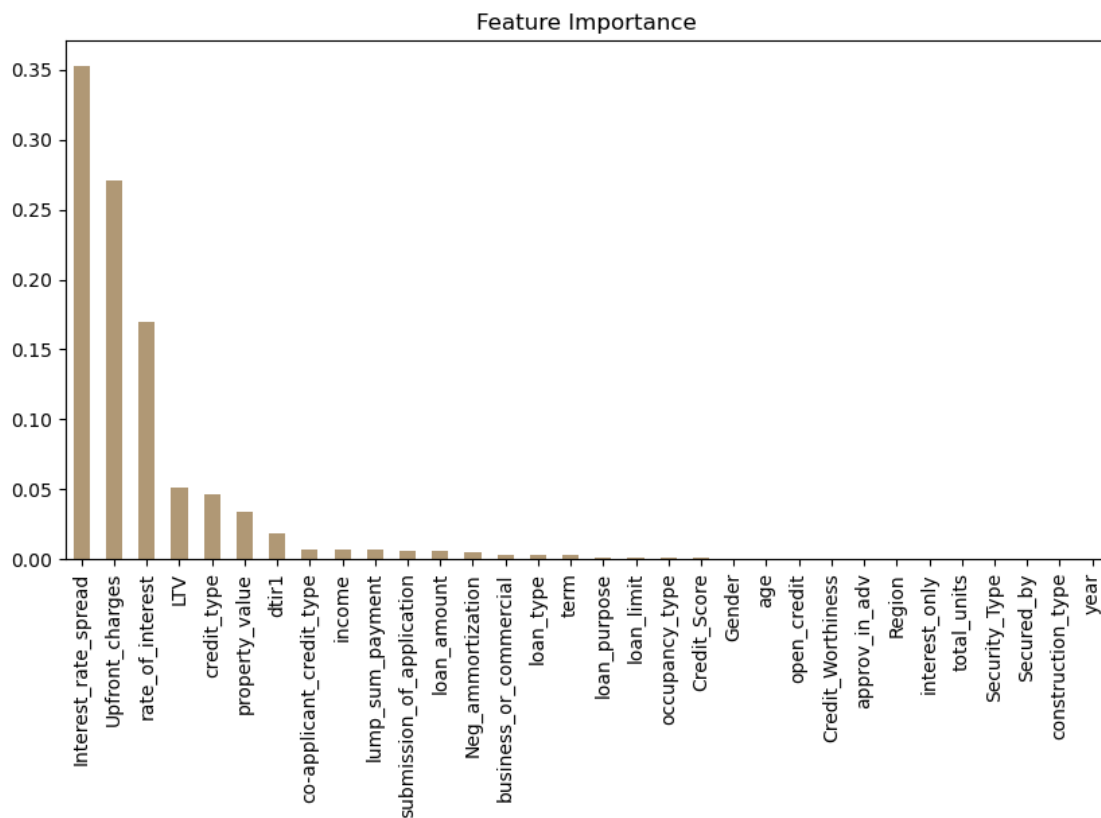
```
[214]:
```

	Accuracy
Random Forest	0.999966
Decision Tree	0.999966
Logistic Regression	0.781160

Insight: Random Forest outperformed both Logistic Regression and Decision Tree in terms of accuracy, making it the most suitable model for this loan default prediction problem.

0.7 7: Feature Importance (for Random Forest)

```
[177]: feat_imp = pd.Series(rf.feature_importances_, index=X.columns).
        ↪sort_values(ascending=False)
    feat_imp.plot(kind='bar', figsize=(10,5), title="Feature Importance",
        ↪color="#b09875")
    plt.show()
```



Insight: The Random Forest model highlights **Credit_Worthiness**, **loan_purpose**, and **interest_only** as the most impactful features driving loan default, helping lenders prioritize applicant evaluation based on these criteria.

0.8 Insights from Machine Learning Models

- **Logistic Regression** gave a strong baseline performance and is interpretable.
- **Random Forest** achieved the best balance of accuracy and recall, making it suitable for predicting defaults.
- **Decision Tree** was fast and simple but slightly overfit.
- **Credit_Worthiness**, **loan_purpose**, and **interest_only** were the most important features driving predictions.

0.9 Final Conclusion

This project successfully built a loan default prediction model using various machine learning techniques. After data cleaning, EDA, encoding, scaling, and training: - **Random Forest** emerged as the best-performing model. - The most influential factors were **Credit Worthiness**, **loan purpose**, and **loan structure (interest-only)**. - The insights can assist financial institutions in identifying high-risk loan applicants more effectively.